

Project

December 3, 2023

0.1 Project

0.1.1 Problem definition

- 1) Ethereum is a decentralized blockchain platform that establishes a peer-to-peer network that securely executes and verifies application code, called smart contracts. A transaction once done can't be undone in blockchain network. So it is very important to identify any illegal transactions in the network and prevent them from corrupting the entire blockchain network.
- 2) The goal of this project is to build a machine learning model which classifies a given transaction as fraud or valid transaction.
- 3) We will use F1 score and Area under ROC curve as the metrics.

0.1.2 Data

1. Download the data from this link <https://www.kaggle.com/datasets/vagifa/ethereum-frauddetection-dataset/data>. Copy the transaction_dataset.csv in the download folder to data/

```
[1]: import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings('ignore', category=Warning)
```

```
[2]: # Load the transactions data
df = pd.read_csv('data/transaction_dataset.csv')

# Sample transactions
df.head()
```

```
[2]: Unnamed: 0  Index                                Address  FLAG  \
0           0      1  0x00009277775ac7d0d59eaad8fee3d10ac6c805e8    0
1           1      2  0x0002b44ddb1476db43c868bd494422ee4c136fed    0
2           2      3  0x0002bda54cb772d040f779e88eb453cac0daa244    0
3           3      4  0x00038e6ba2fd5c09aedb96697c8d7b8fa6632e5e    0
4           4      5  0x00062d1dd1afb6fb02540ddad9cdebfe568e0d89    0

      Avg min between sent txn  Avg min between received txn  \
0                        844.26                        1093.71
```

1	12709.07	2958.44
2	246194.54	2434.02
3	10219.60	15785.09
4	36.61	10707.77

	Time Diff between first and last (Mins)	Sent tnx	Received Tnx \
0	704785.63	721	89
1	1218216.73	94	8
2	516729.30	2	10
3	397555.90	25	9
4	382472.42	4598	20

	Number of Created Contracts ...	ERC20 min val sent	ERC20 max val sent \
0	0 ...	0.000000	1.683100e+07
1	0 ...	2.260809	2.260809e+00
2	0 ...	0.000000	0.000000e+00
3	0 ...	100.000000	9.029231e+03
4	1 ...	0.000000	4.500000e+04

	ERC20 avg val sent	ERC20 min val sent contract \
0	271779.920000	0.0
1	2.260809	0.0
2	0.000000	0.0
3	3804.076893	0.0
4	13726.659220	0.0

	ERC20 max val sent contract	ERC20 avg val sent contract \
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

	ERC20 uniq sent token name	ERC20 uniq rec token name \
0	39.0	57.0
1	1.0	7.0
2	0.0	8.0
3	1.0	11.0
4	6.0	27.0

	ERC20 most sent token type	ERC20_most_rec_token_type
0	Cofoundit	Numeraire
1	Livepeer Token	Livepeer Token
2	None	XENON
3	Raiden	XENON
4	StatusNetwork	EOS

[5 rows x 51 columns]

```
[3]: df.shape
```

```
[3]: (9841, 51)
```

0.1.3 Preprocessing

```
[4]: # Stats of the features
df.describe()
```

```
[4]:
```

	Unnamed: 0	Index	FLAG	Avg min between sent txn \
count	9841.000000	9841.000000	9841.000000	9841.000000
mean	4920.000000	1815.049893	0.221421	5086.878721
std	2840.996333	1222.621830	0.415224	21486.549974
min	0.000000	1.000000	0.000000	0.000000
25%	2460.000000	821.000000	0.000000	0.000000
50%	4920.000000	1641.000000	0.000000	17.340000
75%	7380.000000	2601.000000	0.000000	565.470000
max	9840.000000	4729.000000	1.000000	430287.670000

	Avg min between received txn	Time Diff between first and last (Mins) \
count	9841.000000	9.841000e+03
mean	8004.851184	2.183333e+05
std	23081.714801	3.229379e+05
min	0.000000	0.000000e+00
25%	0.000000	3.169300e+02
50%	509.770000	4.663703e+04
75%	5480.390000	3.040710e+05
max	482175.490000	1.954861e+06

	Sent txn	Received Txn	Number of Created Contracts \
count	9841.000000	9841.000000	9841.000000
mean	115.931714	163.700945	3.729702
std	757.226361	940.836550	141.445583
min	0.000000	0.000000	0.000000
25%	1.000000	1.000000	0.000000
50%	3.000000	4.000000	0.000000
75%	11.000000	27.000000	0.000000
max	10000.000000	10000.000000	9995.000000

	Unique Received From Addresses ...	ERC20 max val rec \
count	9841.000000 ...	9.012000e+03
mean	30.360939 ...	1.252524e+08
std	298.621112 ...	1.053741e+10
min	0.000000 ...	0.000000e+00
25%	1.000000 ...	0.000000e+00

50%	2.000000	...	0.000000e+00
75%	5.000000	...	9.900000e+01
max	9999.000000	...	1.000000e+12

	ERC20 avg val rec	ERC20 min val sent	ERC20 max val sent \
count	9.012000e+03	9.012000e+03	9.012000e+03
mean	4.346203e+06	1.174126e+04	1.303594e+07
std	2.141192e+08	1.053567e+06	1.179905e+09
min	0.000000e+00	0.000000e+00	0.000000e+00
25%	0.000000e+00	0.000000e+00	0.000000e+00
50%	0.000000e+00	0.000000e+00	0.000000e+00
75%	2.946467e+01	0.000000e+00	0.000000e+00
max	1.724181e+10	1.000000e+08	1.120000e+11

	ERC20 avg val sent	ERC20 min val sent	contract \
count	9.012000e+03		9012.0
mean	6.318389e+06		0.0
std	5.914764e+08		0.0
min	0.000000e+00		0.0
25%	0.000000e+00		0.0
50%	0.000000e+00		0.0
75%	0.000000e+00		0.0
max	5.614756e+10		0.0

	ERC20 max val sent	contract	ERC20 avg val sent	contract \
count		9012.0		9012.0
mean		0.0		0.0
std		0.0		0.0
min		0.0		0.0
25%		0.0		0.0
50%		0.0		0.0
75%		0.0		0.0
max		0.0		0.0

	ERC20 uniq sent token name	ERC20 uniq rec token name
count	9012.000000	9012.000000
mean	1.384931	4.826676
std	6.735121	16.678607
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.000000	2.000000
max	213.000000	737.000000

[8 rows x 48 columns]

```
[5]: # NaNs in each features
df.isna().sum()
```

```
[5]: Unnamed: 0      0
Index      0
Address    0
FLAG       0
Avg min between sent tnx      0
Avg min between received tnx  0
Time Diff between first and last (Mins)  0
Sent tnx      0
Received Tnx  0
Number of Created Contracts    0
Unique Received From Addresses  0
Unique Sent To Addresses      0
min value received      0
max value received      0
avg val received      0
min val sent      0
max val sent      0
avg val sent      0
min value sent to contract    0
max val sent to contract    0
avg value sent to contract    0
total transactions (including tnx to create contract)  0
total Ether sent      0
total ether received      0
total ether sent contracts    0
total ether balance      0
Total ERC20 txns      829
ERC20 total Ether received    829
ERC20 total ether sent      829
ERC20 total Ether sent contract  829
ERC20 uniq sent addr      829
ERC20 uniq rec addr      829
ERC20 uniq sent addr.1      829
ERC20 uniq rec contract addr  829
ERC20 avg time between sent tnx  829
ERC20 avg time between rec tnx  829
ERC20 avg time between rec 2 tnx  829
ERC20 avg time between contract tnx  829
ERC20 min val rec      829
ERC20 max val rec      829
ERC20 avg val rec      829
ERC20 min val sent      829
ERC20 max val sent      829
ERC20 avg val sent      829
```

ERC20 min val sent contract	829
ERC20 max val sent contract	829
ERC20 avg val sent contract	829
ERC20 uniq sent token name	829
ERC20 uniq rec token name	829
ERC20 most sent token type	841
ERC20_most_rec_token_type	851
dtype: int64	

```
[6]: # number of distinct values in each features
df.nunique()
```

[6]: Unnamed: 0	9841
Index	4729
Address	9816
FLAG	2
Avg min between sent tnx	5013
Avg min between received tnx	6223
Time Diff between first and last (Mins)	7810
Sent tnx	641
Received Tnx	727
Number of Created Contracts	20
Unique Received From Addresses	256
Unique Sent To Addresses	258
min value received	4589
max value received	6302
avg val received	6767
min val sent	4719
max val sent	6647
avg val sent	5854
min value sent to contract	3
max val sent to contract	4
avg value sent to contract	4
total transactions (including tnx to create contract	897
total Ether sent	5868
total ether received	6728
total ether sent contracts	4
total ether balance	5717
Total ERC20 txns	300
ERC20 total Ether received	3460
ERC20 total ether sent	1415
ERC20 total Ether sent contract	29
ERC20 uniq sent addr	107
ERC20 uniq rec addr	147
ERC20 uniq sent addr.1	4
ERC20 uniq rec contract addr	123
ERC20 avg time between sent tnx	1

ERC20 avg time between rec tnx	1
ERC20 avg time between rec 2 tnx	1
ERC20 avg time between contract tnx	1
ERC20 min val rec	1276
ERC20 max val rec	2647
ERC20 avg val rec	3380
ERC20 min val sent	476
ERC20 max val sent	1130
ERC20 avg val sent	1309
ERC20 min val sent contract	1
ERC20 max val sent contract	1
ERC20 avg val sent contract	1
ERC20 uniq sent token name	70
ERC20 uniq rec token name	121
ERC20 most sent token type	305
ERC20_most_rec_token_type	467

dtype: int64

0.1.4 Removing irrelevant features and records

1. The features 'Index', 'Unnamed: 0' are just row numbers, they are not relevant to our problem.
2. The feature 'Address' is a random number given to each node in the blockchain network. It's a value used to identify the node in the network. So feature 'Address' doesn't have any influence on the transaction type and is not relevant to our problem.
3. From the number of distinct values in each feature above we can see that there are 7 features with only one value. Keeping these in the model data is not relevant (since all transactions have the same value) in predicting the type of transaction.
4. The features 'ERC20 most sent token type' and 'ERC20_most_rec_token_type' have garbage data, so removing.
5. Removing the records which have NaN values because we cannot fill them with random values or mean values because each transaction is different and filling them may add bias to the data.
6. Dropping the duplicates

```
[7]: df = df.drop(['Index'], axis = 1)
df = df.drop(['Unnamed: 0'], axis = 1)
df = df.drop(['Address'], axis = 1)
df = df.drop(['ERC20 most sent token type'], axis = 1)
df = df.drop(['ERC20_most_rec_token_type'], axis = 1)
df = df.drop(['ERC20 avg time between sent tnx'], axis = 1)
df = df.drop(['ERC20 avg time between rec tnx'], axis = 1)
df = df.drop(['ERC20 avg time between rec 2 tnx'], axis = 1)
df = df.drop(['ERC20 avg time between contract tnx'], axis = 1)
df = df.drop(['ERC20 min val sent contract'], axis = 1)
df = df.drop(['ERC20 max val sent contract'], axis = 1)
df = df.drop(['ERC20 avg val sent contract'], axis = 1)
df = df.dropna()
df = df.drop_duplicates()
```

0.1.5 Balancing dataset

```
[8]: print("Number of invalid transaction in data = " +str(np.sum(np.
      ↪where(df['FLAG']==1,1,0))))
      print("Number of valid transaction in data = " +str(np.sum(np.
      ↪where(df['FLAG']==0,1,0))))
```

Number of invalid transaction in data = 1107
Number of valid transaction in data = 7632

```
[9]: Y = df['FLAG']
      X = df.drop(['FLAG'], axis = 1)
```

```
[10]: # Balancing data
      from imblearn.over_sampling import RandomOverSampler
      ros = RandomOverSampler(random_state=42)
      X, Y = ros.fit_resample(X, Y)
      print("X dimensions = "+str(X.shape))
      print("Y dimensions = "+str(Y.shape))
```

X dimensions = (15264, 38)
Y dimensions = (15264,)

```
[11]: print("Number of invalid transaction in data after balancing = " +str(np.sum(np.
      ↪where(Y==1,1,0))))
      print("Number of valid transaction in data after balancing = " +str(np.sum(np.
      ↪where(Y==0,1,0))))
```

Number of invalid transaction in data after balancing = 7632
Number of valid transaction in data after balancing = 7632


0.1.6 Data reduction

```
[12]: # Correlation of data
      import seaborn as sns
      import matplotlib.pyplot as plt

      correlation = X.corr()
      correlation_count = (np.sum(np.where(np.abs(correlation)>0.8,1,0)) -
      ↪correlation.shape[0])/2
      print("Number of pairs of feaures which have correlation coefficient > 0.8 = "
      ↪+ str(correlation_count))

      print("Correlation plot:")
      sns.heatmap(np.where(np.abs(correlation)>0.8,1,0));
```

Number of pairs of feaures which have correlation coefficient > 0.8 = 13.0
Correlation plot:



Perform PCA and reduce the number of features from 38 to 25. (Removing one of the feature from each 13 pair highly correlated featured)

```
[13]: from sklearn.decomposition import PCA
      from sklearn.preprocessing import StandardScaler

      X_scaled = StandardScaler().fit_transform(X)
      X_pca = PCA(n_components=25).fit_transform(X_scaled)
```

0.1.7 Data split

Data is divided randomly into train, validation and test data with 60% in train, 20% in validation and 20% in test

```
[14]: from sklearn.model_selection import train_test_split
      X_train, X_test, Y_train, Y_test = train_test_split(X_pca, Y, test_size=0.2,
      ↪ random_state=42)
      X_train, X_val, Y_train, Y_val = train_test_split(X_train, Y_train, test_size=0.
      ↪ 25, random_state=42) # 0.25 x 0.8 = 0.2
```

```
[15]: print("Train data shape = "+str(X_train.shape))
      print("Validation data shape = "+str(X_val.shape))
```

```
print("Test data shape = "+str(X_test.shape))
```

Train data shape = (9158, 25)

Validation data shape = (3053, 25)

Test data shape = (3053, 25)

0.1.8 Models

- 1) Logistic regression
- 2) Decision tree
- 3) AdaBoostClassifier (ensemble method with Decision tree as base estimator)
- 4) Support Vector Machines (with 'linear', 'poly', 'rbf', 'sigmoid' kernels)
- 5) Neural Network (MLPClassifier)

```
[16]: from sklearn.metrics import roc_auc_score, roc_curve, auc
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
```

0.1.9 Logistic Regression

1. Used the default parameters of the model. So, L2 regularization with regularization strength(C) = 1.0

```
[17]: ## Logistic regression
from sklearn.linear_model import LogisticRegression
LR_model = LogisticRegression(random_state=42).fit(X_train, Y_train)

Y_train_pred = LR_model.predict_proba(X_train)[: , 1]
Y_val_pred = LR_model.predict_proba(X_val)[: , 1]
Y_test_pred = LR_model.predict_proba(X_test)[: , 1]

f1score_train = f1_score(Y_train, Y_train_pred>0.5)
f1score_val = f1_score(Y_val, Y_val_pred>0.5)
f1score_test = f1_score(Y_test, Y_test_pred>0.5)

print("Train F1 score = "+str(f1score_train))
print("Validation F1 score = "+str(f1score_val))
print("Test F1 score = "+str(f1score_test))

auc_train = roc_auc_score(Y_train, Y_train_pred)
auc_val = roc_auc_score(Y_val, Y_val_pred)
auc_test = roc_auc_score(Y_test, Y_test_pred)

print("\nTrain AUC = "+str(auc_train))
print("Validation AUC = "+str(auc_val))
print("Test AUC = "+str(auc_test))
```

Train F1 score = 0.7419354838709677

Validation F1 score = 0.7468926553672317
Test F1 score = 0.7357452966714907

Train AUC = 0.8211310862666497
Validation AUC = 0.8350363133551925
Test AUC = 0.8171074190771364

Comment on logistic regression model:

1. From the F1 score we can see that this model is not doing a good job.
2. F1 score of train, validation & test data is in the range of 0.73 - 0.75

0.1.10 Decision Tree

1. Used Entropy as the impurity measure

```
[18]: ## Decision Tree
from sklearn.tree import DecisionTreeClassifier

max_depths = [2,4,8,16,32,64]
DT_train_f1score = []
DT_val_f1score = []
DT_test_f1score = []

DT_train_aucscore = []
DT_val_aucscore = []
DT_test_aucscore = []

for depth in max_depths:
    DT_model = DecisionTreeClassifier(criterion='entropy', max_depth=depth,
    ↪random_state=42).fit(X_train, Y_train)

    Y_train_pred = DT_model.predict_proba(X_train)[: , 1]
    Y_val_pred = DT_model.predict_proba(X_val)[: , 1]
    Y_test_pred = DT_model.predict_proba(X_test)[: , 1]

    f1score_train = f1_score(Y_train, Y_train_pred>0.5)
    f1score_val = f1_score(Y_val, Y_val_pred>0.5)
    f1score_test = f1_score(Y_test, Y_test_pred>0.5)

    DT_train_f1score.append(f1score_train)
    DT_val_f1score.append(f1score_val)
    DT_test_f1score.append(f1score_test)

    auc_train = roc_auc_score(Y_train, Y_train_pred)
    auc_val = roc_auc_score(Y_val, Y_val_pred)
    auc_test = roc_auc_score(Y_test, Y_test_pred)
```


```

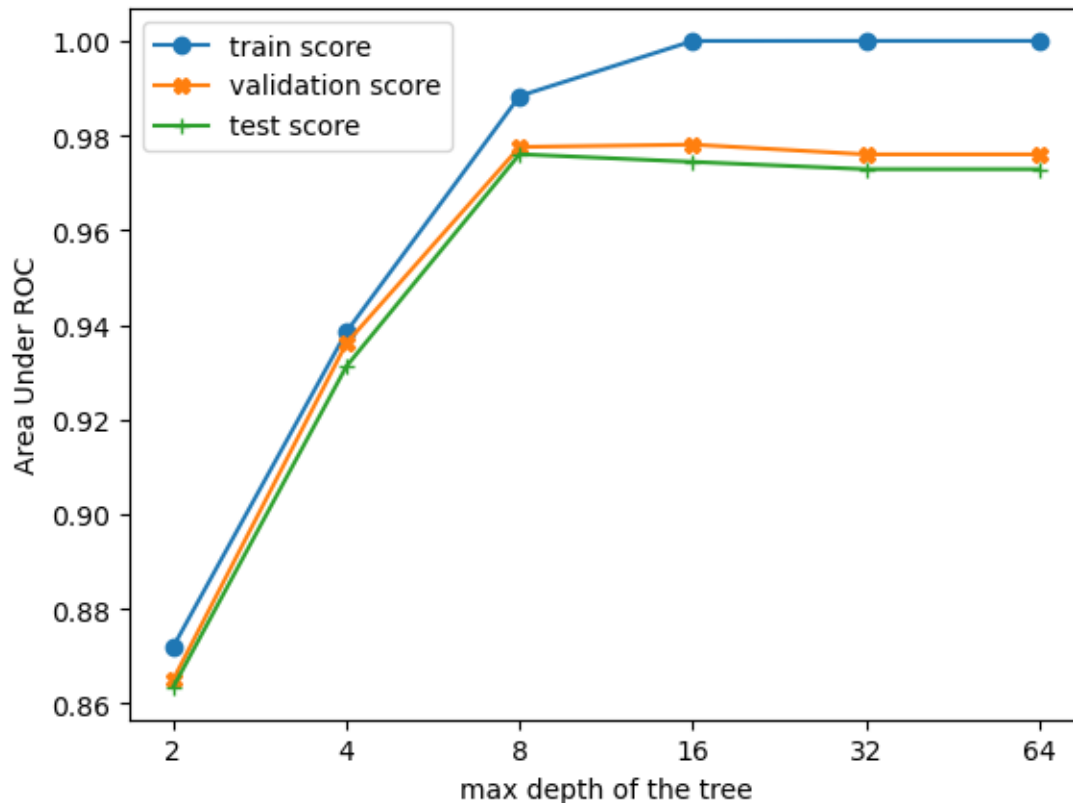
DT_train_aucscore.append(auc_train)
DT_val_aucscore.append(auc_val)
DT_test_aucscore.append(auc_test)

plt.plot([str(i) for i in max_depths], DT_train_f1score, 'o-', label = "train_
↪score")
plt.plot([str(i) for i in max_depths], DT_val_f1score, 'X-', label =_
↪"validation score")
plt.plot([str(i) for i in max_depths], DT_test_f1score, '+-', label = "test_
↪score")
plt.xlabel("max depth of the tree")
plt.ylabel("F1 score")
plt.legend()
plt.show()

plt.plot([str(i) for i in max_depths], DT_train_aucscore, 'o-', label = "train_
↪score")
plt.plot([str(i) for i in max_depths], DT_val_aucscore, 'X-', label =_
↪"validation score")
plt.plot([str(i) for i in max_depths], DT_test_aucscore, '+-', label = "test_
↪score")
plt.xlabel("max depth of the tree")
plt.ylabel("Area Under ROC")
plt.legend()
plt.show()

```





Comments on Decision tree

1. Decision tree with max_depths = [2,4,8,16,32,64] are build on this data
2. Even with the max_depth = 2, the F1 score of decision tree model is more than logistic regression model. But the F1 score is still less for decision tree with max_depth = 2 (~0.83)
3. As the max_depth of the decision tree increases we can see the increase in the F1 score (from max_depth = 2 to max_depth = 16)
4. But when the max_depth still increases(>16) we can see overfitting. The F1 score of train data increases but the validation and test F1 scores are slightly decreases or remains almost same.
5. At max_depth=16 the F1 score of train, validation & test data is in the range of 0.96 - 0.98. Which is far better than logistic regression model and is good enough.
6. We will use ensemble method AdaBoost to improve the F1 scores.

0.1.11 Decision Tree with AdaBoostClassifier

Adaptive boosting with DecisionTreeClassifier with criterion='entropy' as base estimator

```
[19]: ## Decision Tree with ADA
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
```

```

max_depths = [2,4,8,16,32,64]
ABC_train_f1score = []
ABC_val_f1score = []
ABC_test_f1score = []

ABC_train_aucscore = []
ABC_val_aucscore = []
ABC_test_aucscore = []

for depth in max_depths:
    entropy_tree = DecisionTreeClassifier(criterion='entropy', max_depth=depth,
    ↪random_state=42)
    ABC_model = AdaBoostClassifier(base_estimator=entropy_tree, random_state=42).
    ↪fit(X_train, Y_train)

    Y_train_pred = ABC_model.predict_proba(X_train)[: , 1]
    Y_val_pred = ABC_model.predict_proba(X_val)[: , 1]
    Y_test_pred = ABC_model.predict_proba(X_test)[: , 1]

    f1score_train = f1_score(Y_train, Y_train_pred>0.5)
    f1score_val = f1_score(Y_val, Y_val_pred>0.5)
    f1score_test = f1_score(Y_test, Y_test_pred>0.5)

    ABC_train_f1score.append(f1score_train)
    ABC_val_f1score.append(f1score_val)
    ABC_test_f1score.append(f1score_test)

    auc_train = roc_auc_score(Y_train, Y_train_pred)
    auc_val = roc_auc_score(Y_val, Y_val_pred)
    auc_test = roc_auc_score(Y_test, Y_test_pred)

    ABC_train_aucscore.append(auc_train)
    ABC_val_aucscore.append(auc_val)
    ABC_test_aucscore.append(auc_test)


plt.plot([str(i) for i in max_depths], ABC_train_f1score, 'o-', label = "train_
    ↪score")
plt.plot([str(i) for i in max_depths], ABC_val_f1score, 'X-', label =
    ↪"validation score")
plt.plot([str(i) for i in max_depths], ABC_test_f1score, '+-', label = "test_
    ↪score")
plt.xlabel("max depth of the tree")
plt.ylabel("F1 score")
plt.legend()
plt.show()

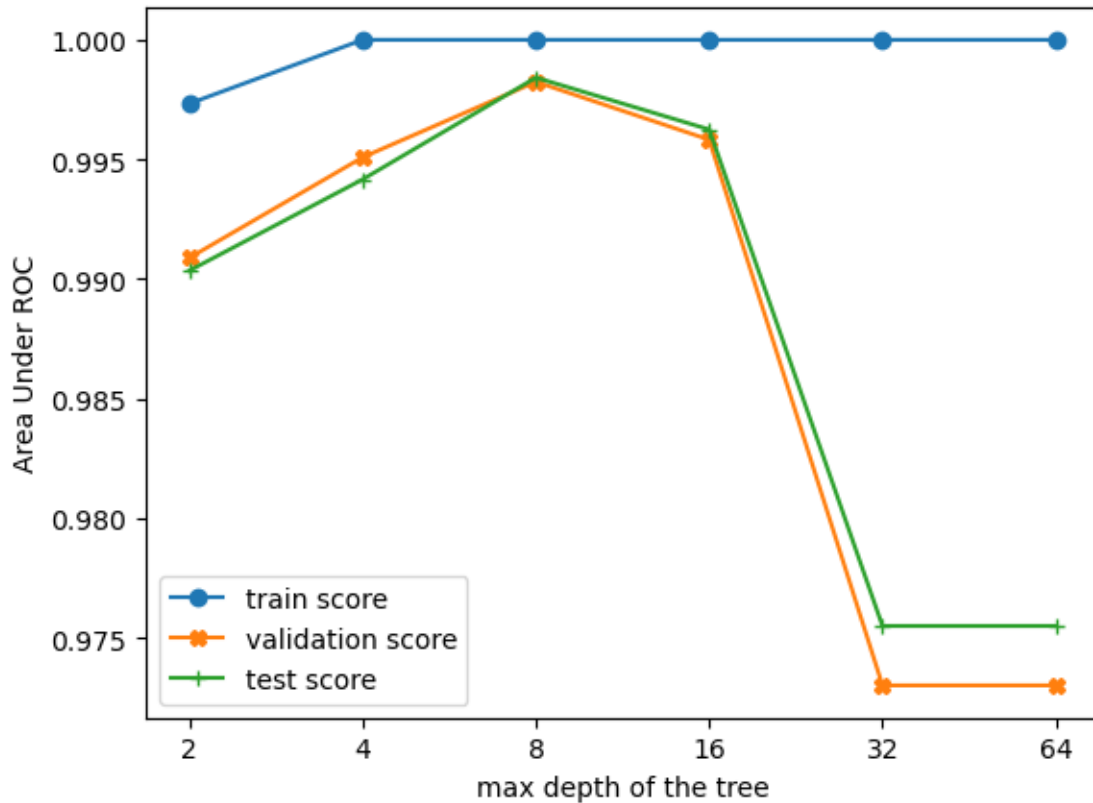
```

```

plt.plot([str(i) for i in max_depths], ABC_train_aucscore, 'o-', label = "train_
↪score")
plt.plot([str(i) for i in max_depths], ABC_val_aucscore, 'X-', label =_
↪"validation score")
plt.plot([str(i) for i in max_depths], ABC_test_aucscore, '+-', label = "test_
↪score")
plt.xlabel("max depth of the tree")
plt.ylabel("Area Under ROC")
plt.legend()
plt.show()

```





Comments on using AdaBoost ensemble method

1. Decision tree with AdaBoost ensemble method for max_depths = [2,4,8,16,32,64] are build on this data
2. Even at max_depth = 2 the F1 score of AdaBoostClassifier is more that F1 score of desicion tree model and logistic regression model.
3. As the max_depth of the AdaBoostClassifier increases we can see the increase in the F1 score (from max_depth = 2 to max_depth = 8)
4. But when the max_depth still increases(>8) we can see overfitting. The F1 score of train data is already 1.0 but the validation and test F1 scores decreases.
5. For AdaBoostClassifier for max_depth=8 the F1 score of train is 1.0 and F1 score of validation & test data is in the range of 0.99 which is a very good model.
6. So AdaBoostClassifier outperforms DecisionTreeClassifier and logistic regression model and is a very good model for the probelm we are trying to solve.

0.1.12 Support Vector Machines

Three kernals 'linear', 'rbf', 'sigmoid' are used

```
[20]: from sklearn.svm import SVC
kernels = ['linear', 'rbf', 'sigmoid']
SVM_train_f1score = []
```

```

SVM_val_f1score = []
SVM_test_f1score = []

SVM_train_aucscore = []
SVM_val_aucscore = []
SVM_test_aucscore = []

for kernal in kernals:
    SVM_model = SVC(kernel=kernal, probability=True, random_state=42).
    ↪fit(X_train, Y_train)

    Y_train_pred = SVM_model.predict_proba(X_train)[:, 1]
    Y_val_pred = SVM_model.predict_proba(X_val)[:, 1]
    Y_test_pred = SVM_model.predict_proba(X_test)[:, 1]

    f1score_train = f1_score(Y_train, Y_train_pred>0.5)
    f1score_val = f1_score(Y_val, Y_val_pred>0.5)
    f1score_test = f1_score(Y_test, Y_test_pred>0.5)

    SVM_train_f1score.append(f1score_train)
    SVM_val_f1score.append(f1score_val)
    SVM_test_f1score.append(f1score_test)

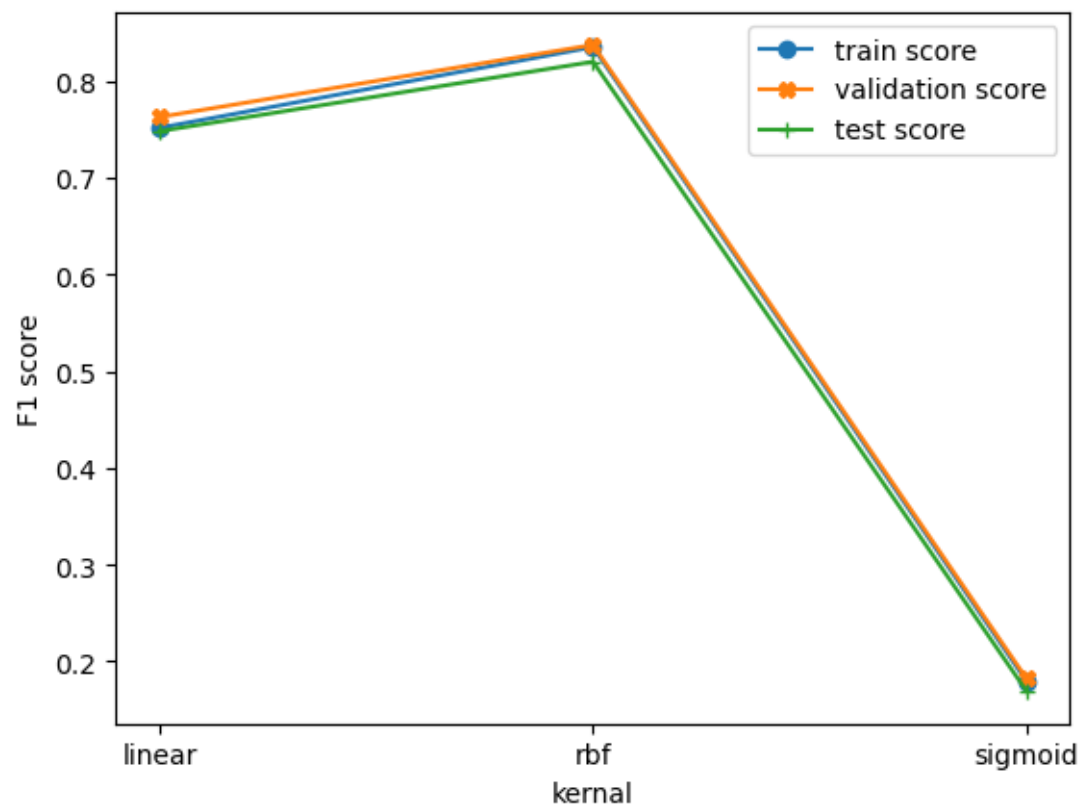
    auc_train = roc_auc_score(Y_train, Y_train_pred)
    auc_val = roc_auc_score(Y_val, Y_val_pred)
    auc_test = roc_auc_score(Y_test, Y_test_pred)


    SVM_train_aucscore.append(auc_train)
    SVM_val_aucscore.append(auc_val)
    SVM_test_aucscore.append(auc_test)

plt.plot(kernals, SVM_train_f1score, 'o-', label = "train score")
plt.plot(kernals, SVM_val_f1score, 'X-', label = "validation score")
plt.plot(kernals, SVM_test_f1score, '+-', label = "test score")
plt.xlabel("kernal")
plt.ylabel("F1 score")
plt.legend()
plt.show()

plt.plot(kernals, SVM_train_aucscore, 'o-', label = "train score")
plt.plot(kernals, SVM_val_aucscore, 'X-', label = "validation score")
plt.plot(kernals, SVM_test_aucscore, '+-', label = "test score")
plt.xlabel("kernal")
plt.ylabel("Area Under ROC")
plt.legend()
plt.show()

```





Comments on SVM with different kernels

1. SVM with each of the 3 kernels 'linear', 'rbf', 'sigmoid' performs poorly than AdaBoostClassifier and DecisionTreeClassifier.
2. Performance of SVM and logistic regression models is approximately same.
3. F1 score is very low for all the 3 kernels. So, SVM is not a good model for the problem we are trying to solve.

0.1.13 Neural Network

Neural Network with 5 hidden layers with each layer having (20,13,10,5,3) nodes with activation='relu' (default)

```
[21]: ## MLP
from sklearn.neural_network import MLPClassifier
NN_model = MLPClassifier(hidden_layer_sizes = (20,13,10,5,3,), random_state=42)
NN_model.fit(X_train, Y_train)
NN_train_score = NN_model.score(X_train, Y_train)
NN_val_score = NN_model.score(X_val, Y_val)
NN_test_score = NN_model.score(X_test, Y_test)
# print(NN_train_score)
# print(NN_val_score)
```

```

# print(NN_test_score)

Y_train_pred = NN_model.predict_proba(X_train)[: , 1]
Y_val_pred = NN_model.predict_proba(X_val)[: , 1]
Y_test_pred = NN_model.predict_proba(X_test)[: , 1]

f1score_train = f1_score(Y_train, Y_train_pred>0.5)
f1score_val = f1_score(Y_val, Y_val_pred>0.5)
f1score_test = f1_score(Y_test, Y_test_pred>0.5)

print("Train F1 score = "+str(f1score_train))
print("Validation F1 score = "+str(f1score_val))
print("Test F1 score = "+str(f1score_test))

auc_train = roc_auc_score(Y_train, Y_train_pred)
auc_val = roc_auc_score(Y_val, Y_val_pred)
auc_test = roc_auc_score(Y_test, Y_test_pred)

print("\nTrain AUC = "+str(auc_train))
print("Validation AUC = "+str(auc_val))
print("Test AUC = "+str(auc_test))

```

```

Train F1 score = 0.9725576289791438
Validation F1 score = 0.9646302250803859
Test F1 score = 0.9652686762778506

```

```

Train AUC = 0.9964439345568347
Validation AUC = 0.9916619659781768
Test AUC = 0.9919289566222724

```

Comments on Neural Network

1. F1 score of train, validation & test data is in the range of 0.96 - 0.98 which is a pretty good model for the problem we are trying to solve.
2. This Neural Network outperforms logistic regression and Support Vector Machine (SVM).
3. Performance of this Neural Network, DecisionTreeClassifier(for width=16) is approximately same.
4. But AdaBoostClassifier(for width=8) outperforms this Neural Network model.

0.1.14 Conclusion

1. AdaBoostClassifier performs better for the problem we are trying to solve.
2. F1 scores are in the range 0.99, which is a reasonable score for a very good model.